

PRN: 123B1B054

Assignment no.: 4

Logistic Regression

1. Logistic Regression is a supervised machine learning algorithm used for classification problems.
2. Unlike Linear Regression (which predicts continuous values), Logistic Regression predicts categorical outcomes (e.g., Yes/No, Spam/Not Spam, Disease/No Disease).
3. The outcome variable is usually binary (0 or 1), but it can also be extended to multiclass classification.
4. Even though it's used for classification, the algorithm is called regression because: It uses a linear combination of input features (just like linear regression).
5. Then applies a logistic (sigmoid) function to map results into probabilities between 0 and 1.

◆ Mathematical Form

For input features

$$X = (x_1, x_2, \dots, x_n)$$

$$z = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

Instead of directly using z, logistic regression applies the sigmoid function:

$$P(Y=1|X) = \sigma(z) = \frac{1}{1+e^{-z}}$$

If $P > 0.5 \rightarrow$ predict class 1 IF $P \leq 0.5 \rightarrow$ predict class 0

◆ Sigmoid (Logistic) Function

1. Converts any real number into a value between 0 and 1.
2. Shape: S-shaped curve.
3. Useful for probability interpretation.

◆ Loss Function

Logistic regression uses Log Loss (Cross-Entropy Loss):

$$L = -\frac{1}{m} \sum [y_i \log(y_i^{\hat{}}) + (1-y_i) \log(1-y_i^{\hat{}})]$$

Ensures model learns to maximize the probability of correct predictions.

◆ Types of Logistic Regression

1. Binary Logistic Regression → Two outcomes (e.g., Pass/Fail).
2. Multinomial Logistic Regression → More than two categories (e.g., Red/Blue/Green).
3. Ordinal Logistic Regression → Ordered categories (e.g., Low/Medium/High).

Assumptions:

1. Dependent variable is categorical.
2. Observations are independent.
3. Predictors should not be highly correlated.

```
import pandas as pd
import numpy as np

df = pd.read_csv("/content/framingham.csv")
df.head(2)
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	I
0	1	39	4.0	0	0.0	0.0	0	0
1	0	46	2.0	0	0.0	0.0	0	0

```
df.duplicated().sum()
```

```
np.int64(0)
```

```
df.isnull().sum()
```

	0
male	0
age	0
education	105
currentSmoker	0
cigsPerDay	29
BPMeds	53
prevalentStroke	0
prevalentHyp	0
diabetes	0
totChol	50
sysBP	0
diaBP	0
BMI	19
heartRate	1
glucose	388
TenYearCHD	0

dtype: int64

```
mode_cols = ["education", "BPMeds"]
for col in mode_cols:
    df[col] = df[col].fillna(df[col].mode()[0])
```

```
median_cols = ["totChol", "BMI", "heartRate", "glucose", "cigsPerDay"]

for col in median_cols:
    df[col] = pd.to_numeric(df[col], errors='coerce') # ensure numeric
    if df[col].notna().any():
        df[col] = df[col].fillna(df[col].median())
    else:
        df[col] = df[col].fillna(0) # fallback if column is empty
```

```
df.isnull().sum()
```

	0
male	0
age	0
education	0
currentSmoker	0
cigsPerDay	0
BPMeds	0
prevalentStroke	0
prevalentHyp	0
diabetes	0
totChol	0
sysBP	0
diaBP	0
BMI	0
heartRate	0
glucose	0
TenYearCHD	0

dtype: int64

df.dtypes

	0
male	int64
age	int64
education	float64
currentSmoker	int64
cigsPerDay	float64
BPMeds	float64
prevalentStroke	int64
prevalentHyp	int64
diabetes	int64
totChol	float64
sysBP	float64
diaBP	float64
BMI	float64
heartRate	float64
glucose	float64
TenYearCHD	int64

dtype: object

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

X = df.drop("TenYearCHD", axis=1)      # Independent variables
y = df["TenYearCHD"]                   # Dependent variable (target)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Feature scaling (important for logistic regression)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Create logistic regression model
model = LogisticRegression(max_iter=1000)
```

```
# Train the model
model.fit(X_train_scaled, y_train)
```

▼ LogisticRegression ⓘ ⓘ

```
LogisticRegression(max_iter=1000)
```

```
# Make predictions
y_pred = model.predict(X_test_scaled)
```

```
# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

print("Accuracy:", accuracy)
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", class_report)
```

Accuracy: 0.8573113207547169

Confusion Matrix:

```
[[718  6]
 [115  9]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.86	0.99	0.92	724
1	0.60	0.07	0.13	124
accuracy			0.86	848
macro avg	0.73	0.53	0.53	848
weighted avg	0.82	0.86	0.81	848

Decision Tree Model

a) What is a Decision Tree?

A Decision Tree is a supervised machine learning algorithm used for classification and regression tasks. It works by splitting the dataset into smaller subsets based on feature values, forming a tree-like structure where:

Each internal node represents a decision on a feature.

Each branch represents the outcome of the decision.

Each leaf node represents the final class or prediction.

b) How Does It Work?

The tree is constructed using a top-down, greedy approach:

The algorithm selects the best feature to split the data using metrics like:

Gini Impurity

Entropy (Information Gain)

It continues splitting the data recursively until:

All samples in a node belong to the same class.

A stopping criterion is reached (maximum depth, minimum samples per leaf).

```
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, classification_report

X = df.drop("TenYearCHD", axis=1)
y = df["TenYearCHD"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Create Decision Tree model
dt_model = DecisionTreeClassifier(max_depth=4, random_state=42, class_weight

# Train the model
dt_model.fit(X_train_scaled, y_train)

# Predictions
y_pred = dt_model.predict(X_test_scaled)

# Evaluation
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

print("Accuracy:", accuracy)
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", class_report)
```

Accuracy: 0.6014150943396226

Confusion Matrix:

```
[[429 290]
 [ 48 81]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.90	0.60	0.72	719
1	0.22	0.63	0.32	129
accuracy			0.60	848
macro avg	0.56	0.61	0.52	848
weighted avg	0.80	0.60	0.66	848

```
from sklearn.tree import DecisionTreeClassifier

# Train an initial tree to get feature importances
initial_tree = DecisionTreeClassifier(random_state=42)
initial_tree.fit(X_train_scaled, y_train)

# Get feature importances
importances = initial_tree.feature_importances_
feature_importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

print(feature_importance_df)
```

	Feature	Importance
12	BMI	0.160206
10	sysBP	0.147702
11	diaBP	0.126375
9	totChol	0.117452
1	age	0.114291
14	glucose	0.104362
13	heartRate	0.102665
4	cigsPerDay	0.050882
2	education	0.037796
0	male	0.015689
3	currentSmoker	0.006021
5	BPMeds	0.005318
6	prevalentStroke	0.004721
8	diabetes	0.004364
7	prevalentHyp	0.002157

```
top_features = feature_importance_df['Feature'][:6].tolist()
print("Top features:", top_features)

X_top = X[top_features]

# Split and scale again
X_train_top, X_test_top, y_train, y_test = train_test_split(
    X_top, y, test_size=0.2, random_state=42, stratify=y
)

scaler = StandardScaler()
X_train_scaled_top = scaler.fit_transform(X_train_top)
X_test_scaled_top = scaler.transform(X_test_top)
```

Top features: ['BMI', 'sysBP', 'diaBP', 'totChol', 'age', 'glucose']

```
# Train tree on top features only
dt_model_top = DecisionTreeClassifier(
    max_depth=4, class_weight='balanced', random_state=42, criterion="entropy"
)
dt_model_top.fit(X_train_scaled_top, y_train)
```

```
# Predict and evaluate
y_pred_top = dt_model_top.predict(X_test_scaled_top)

from sklearn.metrics import accuracy_score, confusion_matrix, classification
print("Accuracy:", accuracy_score(y_test, y_pred_top))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_top))
print("Classification Report:\n", classification_report(y_test, y_pred_top))
```

Accuracy: 0.6757075471698113

Confusion Matrix:

```
[[510 209]
 [ 66  63]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.71	0.79	719
1	0.23	0.49	0.31	129
accuracy			0.68	848
macro avg	0.56	0.60	0.55	848
weighted avg	0.79	0.68	0.72	848

```
import matplotlib.pyplot as plt
from sklearn import tree
```

```
plt.figure(figsize=(45,25))
tree.plot_tree(
    dt_model,
    feature_names=df.columns,           # column names
    filled=True,                        # color nodes
    rounded=True,                       # rounded boxes
    fontsize=12                         # readable text
)
plt.show()
```

